

# Using Dynamic SQL

**ORACLE®**

Oracle Academy Study Materials

# Using Dynamic SQL

All SQL statements in the database go through various stages:

1. Parse: Pre-execution “is this possible?” checks syntax, object existence, privileges, and so on
2. Bind: Getting the actual values of any variables referenced in the statement
3. Execute: The statement is executed.
4. Fetch: Results are returned to the user.

Some stages might not be relevant for all statements:

- The fetch phase is applicable to queries.
- For embedded SQL statements such as SELECT, DML, MERGE, COMMIT, SAVEPOINT, and ROLLBACK, the parse and bind phases are done at compile time.
- For dynamic SQL statements, all phases are performed at run time.

# Using Dynamic SQL

- When a SQL statement is included in a PL/SQL subprogram, the parse and bind phases are normally done at compile time, that is, when the procedure or function is created.
- If the text of the SQL statement is not known when the procedure is created, the Oracle server cannot parse it.
- Example:

```
CREATE PROCEDURE drop_any_table(p_table_name
VARCHAR2)
IS BEGIN
    DROP TABLE p_table_name; -- cannot be parsed
END;
```

# Using Dynamic SQL

The embedded SQL statements available in PL/SQL are limited to :

- SELECT, INSERT, UPDATE, DELETE, COMMIT, and ROLLBACK,
- All of which are parsed at compile time; that is, they have a fixed structure.

Use Dynamic SQL functionality if you require:

- The structure of a SQL statement to be altered at execution time
- To access to DDL statements and other SQL functionality in PL/SQL
- To create a SQL statement whose text is not completely known in advance.

Dynamic SQL:

- Is constructed and stored as a character string within a subprogram.
- Is a SQL statement with varying column data, or different conditions with or without bind variables.
- Enables data-definition, data-control, or session-control statements to be written and executed from PL/SQL.

# Using Dynamic SQL

- PL/SQL does not support DDL statements written directly in a program.
- To perform these kinds of tasks in PL/SQL, you must construct SQL statements dynamically in character strings and execute them using either of the following:
  - Native Dynamic SQL statements with EXECUTE IMMEDIATE
  - The DBMS\_SQL package
- The process of using SQL statements that are not embedded in your source program and are constructed in strings and executed at run time is known as “dynamic SQL.”

# Using Dynamic SQL

- Native Dynamic SQL (NDS) allows you to work around this by constructing and storing SQL as a character string within a subprogram.
- NDS:
  - Provides native support for Dynamic SQL directly in the PL/SQL language.
  - Enables data-definition, data-control, or session-control statements to be written and executed from PL/SQL.
  - Provides the ability to execute SQL statements whose structure is unknown until execution time.
  - Can also use the OPEN-FOR, FETCH, and CLOSE PL/SQL statements.

# Using Dynamic SQL

## Using the EXECUTE IMMEDIATE Statement:

```
EXECUTE IMMEDIATE dynamic_string
  [INTO {define_variable
    [, define_variable] ... | record}]
  [USING [IN|OUT|IN OUT] bind_argument
    [, [IN|OUT|IN OUT] bind_argument] ... ];
```

- INTO is used for single-row queries and specifies the variables or records into which column values are retrieved.
- USING holds all bind arguments. The default parameter mode is IN, if not specified.
- *dynamic\_string* is a character variable or literal containing the text of a SQL statement.
- *define\_variable* is a PL/SQL variable that stores a selected column value.
- *record* is a user-defined or %ROWTYPE record that stores a selected row.
- *bind\_argument* is an expression whose value is passed to the dynamic SQL statement at execution time.

# Available Methods for Using NDS

Method #	SQL Statement Type	NDS SQL Statements Used
<b>Method 1</b>	<b>Non-query without</b> host variables	EXECUTE IMMEDIATE without the USING and INTO clauses
<b>Method 2</b>	<b>Non-query with</b> known number of input host variables	EXECUTE IMMEDIATE with a USING clause
<b>Method 3</b>	<b>Query with known</b> number of select-list items and input host variables	EXECUTE IMMEDIATE with the USING and INTO clauses
<b>Method 4</b>	<b>Query with unknown</b> number of select-list items or input host variables	Use the DBMS_SQL package



# Available Methods for Using NDS

## Method 1:

- This method lets the program accept or build a dynamic SQL statement, then immediately execute it using the EXECUTE IMMEDIATE command. The SQL statement must not be a query (SELECT statement) and must not contain any placeholders for input host variables.
- With Method 1, the SQL statement is parsed every time it is executed.
- Examples of non-queries include data definition language (DDLs) statements, UPDATEs, INSERTs, or DELETEs.

# Available Methods for Using NDS

## Method 2:

- This method lets the program accept or build a dynamic SQL statement, then process it using the PREPARE and EXECUTE commands.
- The SQL statement must not be a query.
- The number of placeholders for input host variables and the data types of the input host variables must be known at precompile time.
- Examples:
  - `INSERT INTO EMPLOYEES (FIRST_NAME, LAST_NAME, JOB_ID) VALUES (:emp_first_name, :emp_last_name, :job_id)`
  - `DELETE FROM EMPLOYEES WHERE EMPLOYEE_ID = :emp_number`
- With Method 2, the SQL statement is parsed just once, but can be executed many times with different values for the host variables.
- SQL data definition statements such as CREATE and GRANT are executed when they are prepared.

# Available Methods for Using NDS

## Method 3:

- This method lets the program accept or build a dynamic query, then process it using the PREPARE command with the DECLARE, OPEN, FETCH, and CLOSE cursor commands.
- The number of select-list items, the number of placeholders for input host variables, and the data types of the input host variables must be known at precompile time.
- Examples:
  - ```
SELECT DEPARTMENT_ID, MIN(SALARY), MAX(SALARY)
FROM EMPLOYEES
GROUP BY DEPARTMENT_ID
```
  - ```
SELECT LAST_NAME, EMPLOYEE_ID
FROM EMPLOYEES
WHERE DEPARTMENT_ID = :dept_number
```

# Available Methods for Using NDS

## Method 4:

- This method lets your program accept or build a dynamic SQL statement, then process it using descriptors.
- A descriptor is an area of memory used by your program and Oracle to hold a complete description of the variables in a dynamic SQL statement.
- The number of select-list items, the number of placeholders for input host variables, and the data types of the input host variables can be unknown until run time.
- Examples:
  - `INSERT INTO EMPLOYEES (<unknown>) VALUES (<unknown>)`
  - `SELECT <unknown> FROM EMPLOYEES WHERE DEPARTMENT_ID = 20`
- Method 4 is required for dynamic SQL statements that contain an unknown number of select-list items or input host variables.
- With this method, you use the `DBMS_SQL` package.

# Using Dynamic SQL – Example 1

- Constructing the dynamic statement in-line:

```
CREATE PROCEDURE drop_any_table(p_table_name VARCHAR2) IS
BEGIN
    EXECUTE IMMEDIATE 'DROP TABLE ' || p_table_name;
END;
```

- Constructing the dynamic statement in a variable:

```
CREATE PROCEDURE drop_any_table(p_table_name VARCHAR2) IS
    v_dynamic_stmt VARCHAR2(50);
BEGIN
    v_dynamic_stmt := 'DROP TABLE ' || p_table_name;
    EXECUTE IMMEDIATE v_dynamic_stmt;
END;
```

```
BEGIN drop_any_table('EMPLOYEE_NAMES'); END;
```

# Using Dynamic SQL – Example 2

- Deleting all the rows from any table and returning a count:

```
CREATE FUNCTION del_rows(p_table_name VARCHAR2)
RETURN NUMBER IS
BEGIN
    EXECUTE IMMEDIATE 'DELETE FROM ' || p_table_name;
    RETURN SQL%ROWCOUNT;
END;
```

- Invoking the function:

```
DECLARE
    v_count    NUMBER;
BEGIN
    v_count := del_rows('EMPLOYEE_NAMES');
    DBMS_OUTPUT.PUT_LINE(v_count || ' rows deleted.');
```

```
END;
```

# Using Dynamic SQL – Example 3


- Inserting a row into a table with two columns and invoking the procedure:

```
CREATE PROCEDURE add_row(p_table_name VARCHAR2,  
    p_id NUMBER, p_name VARCHAR2) IS  
BEGIN  
    EXECUTE IMMEDIATE 'INSERT INTO ' || p_table_name ||  
        'VALUES(' || p_id || ', ' || p_name || ')';  
END;
```

```
BEGIN  
    add_row('EMPLOYEE_NAMES', 250, 'Chang');  
END;
```

# Using the DBMS\_SQL Package

- Some of the procedures and functions of the DBMS\_SQL package are:
- OPEN\_CURSOR
- PARSE
- BIND\_VARIABLE
- EXECUTE
- FETCH\_ROWS
- CLOSE\_CURSOR





# Using DBMS\_SQL with a DML Statement

- Example of deleting rows:

```
CREATE OR REPLACE FUNCTION del_rows
(p_table_name VARCHAR2) RETURN NUMBER IS
  v_csr_id  INTEGER;
  v_rows_del  NUMBER;
BEGIN
  v_csr_id := DBMS_SQL.OPEN_CURSOR;
  DBMS_SQL.PARSE(v_csr_id,
    'DELETE FROM ' || p_table_name, DBMS_SQL.NATIVE);
  v_rows_del := DBMS_SQL.EXECUTE(v_csr_id);
  DBMS_SQL.CLOSE_CURSOR(v_csr_id);
  RETURN v_rows_del;
END;
```

```
CREATE FUNCTION del_rows(p_table_name VARCHAR2)
RETURN NUMBER IS
BEGIN
  EXECUTE IMMEDIATE 'DELETE FROM ' || p_table_name;
  RETURN SQL%ROWCOUNT;
END;
```

# Using DBMS\_SQL with a Parameterized DML Statement

```
CREATE PROCEDURE add_row (p_table_name VARCHAR2,  
    p_id NUMBER, p_name VARCHAR2) IS  
    v_csr_id    INTEGER;  
    v_stmt      VARCHAR2(200);  
    v_rows_added NUMBER;  
BEGIN  
    v_stmt := 'INSERT INTO ' || p_table_name ||  
        ' VALUES(' || p_id || ', ' || p_name || ')';  
    v_csr_id := DBMS_SQL.OPEN_CURSOR;  
    DBMS_SQL.PARSE(v_csr_id, v_stmt, DBMS_SQL.NATIVE);  
    v_rows_added := DBMS_SQL.EXECUTE(v_csr_id);  
    DBMS_SQL.CLOSE_CURSOR(v_csr_id);  
END;
```

```
CREATE PROCEDURE add_row(p_table_name VARCHAR2,  
    p_id NUMBER, p_name VARCHAR2) IS  
BEGIN  
    EXECUTE IMMEDIATE 'INSERT INTO ' || p_table_name ||  
        'VALUES(' || p_id || ', ' || p_name || ')';  
END;
```

# Using Dynamic SQL

## Comparison of Native Dynamic SQL and the DBMS\_SQL Package

- Native Dynamic SQL:
  - Is easier to use than DBMS\_SQL
  - Requires less code than DBMS\_SQL
  - Often executes faster than DBMS\_SQL because there are fewer statements to execute.



# Improving PL/SQL Performance

**ORACLE®**

Oracle Academy Study Materials

# Improving PL/SQL Performance

## Purpose:

- Until now, you have learned how to write, compile, and execute PL/SQL code without thinking much about how long the execution will take.
- None of the tables in HR schema contain more than a few hundred rows, so the execution is always fast.
- But in real organizations, tables can contain millions or even billions of rows.
- Obviously, processing two million rows takes much longer than processing twenty rows, and therefore some ways to speed up the processing of very large sets of data should be considered.

# Using the NOCOPY Hint

- In PL/SQL and most other programming languages, there are two ways to pass parameter arguments between a calling program and a called subprogram:
  - by value,
  - by reference.
- Passing by value means that the argument values are copied from the calling program's memory to the subprogram's memory, and copied back again when the subprogram is exited.
  - while the subprogram is executing, there are two copies of each argument.
- Passing by reference means that the argument values are not copied.
  - the two programs share a single copy of the data.
- While passing by value is safer, it can use a lot of memory and execute slowly if the argument value is large.

# Using the NOCOPY Hint

- Example:

```
CREATE OR REPLACE PACKAGE emp_pkg IS
  TYPE t_emp IS TABLE OF employees%ROWTYPE
  INDEX BY BINARY_INTEGER;
  PROCEDURE emp_proc
    (p_small_arg IN NUMBER, p_big_arg OUT t_emp);
  ...
END emp_pkg;
```

- Suppose EMP\_PKG.EMP\_PROC fetches one million EMPLOYEES rows into P\_BIG\_ARG.
- And those one million rows must be copied to the calling environment at the end of the procedure's execution.

# Using the NOCOPY Hint

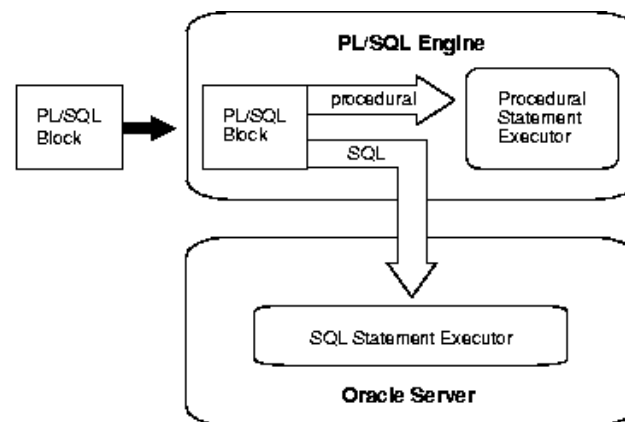
- By default, PL/SQL IN parameter arguments are passed by reference, while OUT and IN OUT arguments are passed by value.
- To pass an OUT or IN OUT argument by reference the NOCOPY hint can be used.

```
CREATE OR REPLACE PACKAGE emp_pkg IS
  TYPE t_emp IS TABLE OF employees%ROWTYPE
  INDEX BY BINARY_INTEGER;
  PROCEDURE emp_proc
    (p_small_arg IN NUMBER, p_big_arg OUT NOCOPY t_emp);
  ...
END emp_pkg;
```



# Bulk Binding

- Many PL/SQL blocks contain both PL/SQL statements and SQL statements, each of which is executed by a different part of the Oracle software called the PL/SQL Engine and the SQL Engine.
- A change from one engine to the other is called a context switch, which takes time. For one change, this is at most a few milliseconds, but what if there are millions of changes?
- If we FETCH (in a cursor) and process millions of rows one at a time, that's millions of context switches.
- FETCH is a SQL statement because it accesses database tables, but the processing is done by PL/SQL statements.



# Bulk Binding

```
CREATE OR REPLACE PROCEDURE fetch_all_emps IS
  CURSOR emp_curs IS SELECT * FROM employees;
BEGIN
  FOR v_emprec IN emp_curs LOOP
    DBMS_OUTPUT.PUT_LINE(v_emprec.first_name);
  END LOOP;
END fetch_all_emps;
```

- It would be much quicker to fetch all the rows in just one context switch within the SQL Engine - this is what Bulk Binding does.
- But: If each row is (on average) 100 bytes in size, storing one million rows will need 100 megabytes of memory.
- So Bulk Binding is a trade-off: more memory required (possibly bad) but faster execution (good).

# Bulk Binding

## Bulk Binding a SELECT: Using BULK COLLECT:

```
CREATE OR REPLACE PROCEDURE fetch_all_emps IS
    TYPE t_emp IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
    v_emptab t_emp;
BEGIN
    SELECT * BULK COLLECT INTO v_emptab FROM employees;
    FOR i IN v_emptab.FIRST..v_emptab.LAST LOOP
        IF v_emptab.EXISTS(i) THEN
            DBMS_OUTPUT.PUT_LINE(v_emptab(i).last_name);
        END IF;
    END LOOP;
END fetch_all_emps;
```

- When using BULK COLLECT, we do not declare a cursor because we do not fetch individual rows one at a time.
- Instead, we SELECT the whole database table into the PL/SQL INDEX BY table in a single SQL statement.

# Bulk Binding

## Bulk Binding a SELECT: Using BULK COLLECT – Example 2:

```
CREATE OR REPLACE PROCEDURE fetch_some_emps IS
    TYPE t_salary IS TABLE OF employees.salary%TYPE
        INDEX BY BINARY_INTEGER;
    v_saltab t_salary;
BEGIN
    SELECT salary BULK COLLECT INTO v_saltab
    FROM employees WHERE department_id = 20 ORDER BY salary;
    FOR i IN v_saltab.FIRST..v_saltab.LAST LOOP
        IF v_saltab.EXISTS(i) THEN
            DBMS_OUTPUT.PUT_LINE(v_saltab(i));
        END IF;
    END LOOP;
END fetch_some_emps;
```

# Bulk Binding

```
CREATE OR REPLACE PROCEDURE insert_emps IS
    TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
    v_emptab t_emps;
BEGIN
    FOR i IN v_emptab.FIRST..v_emptab.LAST LOOP
        INSERT INTO employees VALUES v_emptab(i);
    END LOOP;
END insert_emps;
```

- Bulk Binding with DML: Using FORALL (the code will compile, but will not perform any inserts as the v\_emptab table is not populated).

```
CREATE OR REPLACE PROCEDURE insert_emps IS
    TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
    v_emptab t_emps;
BEGIN
    FORALL i IN v_emptab.FIRST..v_emptab.LAST
        INSERT INTO employees VALUES v_emptab(i);
END insert_emps;
```

# Bulk Binding

- We can combine BULK COLLECT and FORALL:

```
CREATE OR REPLACE PROCEDURE copy_emps IS
    TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
    v_emptab t_emps;
BEGIN
    SELECT * BULK COLLECT INTO v_emptab FROM employees;
    FORALL i IN v_emptab.FIRST..v_emptab.LAST
        INSERT INTO new_employees VALUES v_emptab(i);
END copy_emps;
```

- Since no columns are specified in the INSERT statement, the record structure of the collection must match the table exactly.
- Bulk binds can also improve the performance when loading collections from queries.

# Bulk Binding

- We can use FORALL with UPDATE and DELETE statements as well as with INSERT:

```
CREATE OR REPLACE PROCEDURE copy_emps IS
  TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
  v_emptab t_emps;
BEGIN
  SELECT * BULK COLLECT INTO v_emptab FROM employees;
  FORALL i IN v_emptab.FIRST..v_emptab.LAST
    INSERT INTO new_employees VALUES v_emptab(i);
END copy_emps;
```

# Bulk Binding

- In addition to implicit cursor attributes such as SQL%ROWCOUNT, Bulk Binding uses two extra cursor attributes, which are both INDEX BY tables.

```
CREATE OR REPLACE PROCEDURE insert_emps IS
  TYPE t_emps IS TABLE OF employees%ROWTYPE
  INDEX BY BINARY_INTEGER;
  v_emptab t_emps;
BEGIN
  SELECT * BULK COLLECT INTO v_emptab FROM employees;
  FORALL i IN v_emptab.FIRST..v_emptab.LAST
    INSERT INTO emp VALUES v_emptab(i);
  FOR i IN v_emptab.FIRST..v_emptab.LAST LOOP
    DBMS_OUTPUT.PUT_LINE('Inserted: '
      || i || ' ' || SQL%BULK_ROWCOUNT(i) || ' rows');
  END LOOP;
END insert_emps;
```

- SQL%BULK\_ROWCOUNT(i) shows the number of rows processed by the i-th execution of a DML statement when using FORALL.



# Bulk Binding

- Bulk Binding Cursor Attributes: SQL%BULK\_EXCEPTIONS – example:

```
CREATE OR REPLACE PROCEDURE insert_emps IS
  TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
  v_emptab t_emps;
BEGIN
  SELECT * BULK COLLECT INTO v_emptab FROM employees;
  FORALL i IN v_emptab.FIRST..v_emptab.LAST
    INSERT INTO employees VALUES v_emptab(i);
END insert_emps;
```

- If one of the INSERTs fails, perhaps because a constraint was violated, the whole FORALL statement fails.
- No rows are inserted and there is no information on which row failed to insert.

# Bulk Binding

- If SAVE EXCEPTIONS is added to FORALL statement:

```
CREATE OR REPLACE PROCEDURE insert_emps IS
  TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
  v_emptab t_emps;
BEGIN
  SELECT * BULK COLLECT INTO v_emptab FROM employees;
  FORALL i IN v_emptab.FIRST..v_emptab.LAST SAVE EXCEPTIONS
    INSERT INTO employees VALUES v_emptab(i);
END insert_emps;
```

- All the non-violating rows will be inserted.
- The violating rows populate an INDEX BY table called SQL%BULK\_EXCEPTIONS which has two fields:
  - ERROR\_INDEX - shows which inserts failed (first, second, ...),
  - ERROR\_CODE – shows the Oracle Server predefined error code.

# Bulk Binding

- SQL%BULK\_EXCEPTIONS in the EXCEPTION section:

```
CREATE OR REPLACE PROCEDURE insert_emps IS
    TYPE t_emps IS TABLE OF employees%ROWTYPE INDEX BY BINARY_INTEGER;
    v_emptab t_emps;
BEGIN
    SELECT * BULK COLLECT INTO v_emptab FROM employees;
    FORALL i IN v_emptab.FIRST..v_emptab.LAST SAVE EXCEPTIONS
        INSERT INTO employees VALUES v_emptab(i);
EXCEPTION
    WHEN OTHERS THEN
        FOR j IN 1..SQL%BULK_EXCEPTIONS.COUNT LOOP
            DBMS_OUTPUT.PUT_LINE(SQL%BULK_EXCEPTIONS(j).ERROR_INDEX);
            DBMS_OUTPUT.PUT_LINE(SQL%BULK_EXCEPTIONS(j).ERROR_CODE);
        END LOOP;
END insert_emps;
```

# Bulk Binding

SQL%BULK\_EXCEPTIONS summary:

- The FORALL statement includes an optional SAVE EXCEPTIONS clause that allows bulk operations to save exception information and continue processing.
- Once the operation is complete, the exception information can be retrieved using the SQL%BULK\_EXCEPTIONS attribute.
- This is a collection of exceptions for the most recently executed FORALL statement, with the following two fields for each exception:
  - SQL%BULK\_EXCEPTIONS(i).ERROR\_INDEX
  - SQL%BULK\_EXCEPTIONS(i).ERROR\_CODE

# Using the RETURNING Clause

- Sometimes we need to perform a DML operation on a row, and then SELECT column values from the updated row for later use.
- Two SQL statements are required: an UPDATE and a SELECT:

```
CREATE OR REPLACE PROCEDURE update_one_emp
  (p_emp_id    IN  employees.employee_id%TYPE,
   p_salary_raise_percent IN  NUMBER) IS
  v_new_salary  employees.salary%TYPE;
BEGIN
  UPDATE employees
  SET salary = salary * (1 + p_salary_raise_percent)
  WHERE employee_id = p_emp_id;
  SELECT salary INTO v_new_salary
  FROM employees
  WHERE employee_id = p_emp_id;
  DBMS_OUTPUT.PUT_LINE('New salary is: ' || v_new_salary);
END update_one_emp;
```

# Using the RETURNING Clause

- The SELECT can be done within the UPDATE statement – it is faster because only one call is made to the SQL Engine:

```
CREATE OR REPLACE PROCEDURE update_one_emp
  (p_emp_id    IN  employees.employee_id%TYPE,
   p_salary_raise_percent IN  NUMBER) IS
  v_new_salary  employees.salary%TYPE;
BEGIN
  UPDATE employees
  SET salary = salary * (1 + p_salary_raise_percent)
  WHERE employee_id = p_emp_id
  RETURNING salary INTO v_new_salary;
  DBMS_OUTPUT.PUT_LINE('New salary is: ' || v_new_salary);
END update_one_emp;
```

# Using the RETURNING Clause with FORALL

- if we want to update millions of rows and see the updated values, we can use RETURNING with a Bulk Binding FORALL clause:

```
CREATE OR REPLACE PROCEDURE update_all_emps
  (p_salary_raise_percent IN  NUMBER) IS
  TYPE t_empid IS TABLE OF employees.employee_id%TYPE
  INDEX BY BINARY_INTEGER;
  TYPE t_sal IS TABLE OF employees.salary%TYPE
  INDEX BY BINARY_INTEGER;
  v_empidtab  t_empid;
  v_salstab  t_sal;
BEGIN
  SELECT employee_id BULK COLLECT INTO v_empidtab FROM employees;
  FORALL i IN v_empidtab.FIRST..v_empidtab.LAST
  UPDATE employees
  SET salary = salary * (1 + p_salary_raise_percent)
  WHERE employee_id = v_empidtab(i)
  RETURNING salary BULK COLLECT INTO v_salstab;
END update_all_emps;
```

# Data Dictionary

**ORACLE®**

Oracle Academy Study Materials



# Data Dictionary

- Imagine that you have created many procedures and/or functions, as well as tables and other database objects.
- The **Data Dictionary** remembers this information for you.
- Every Oracle database contains a Data Dictionary – „an automatically-managed master catalog of everything in the database”.
- All database objects, such as tables, views, users and their privileges, procedures or functions are automatically registered in the Data Dictionary when they are created.
- If an object is later altered or dropped, the Dictionary is automatically updated to reflect the change.

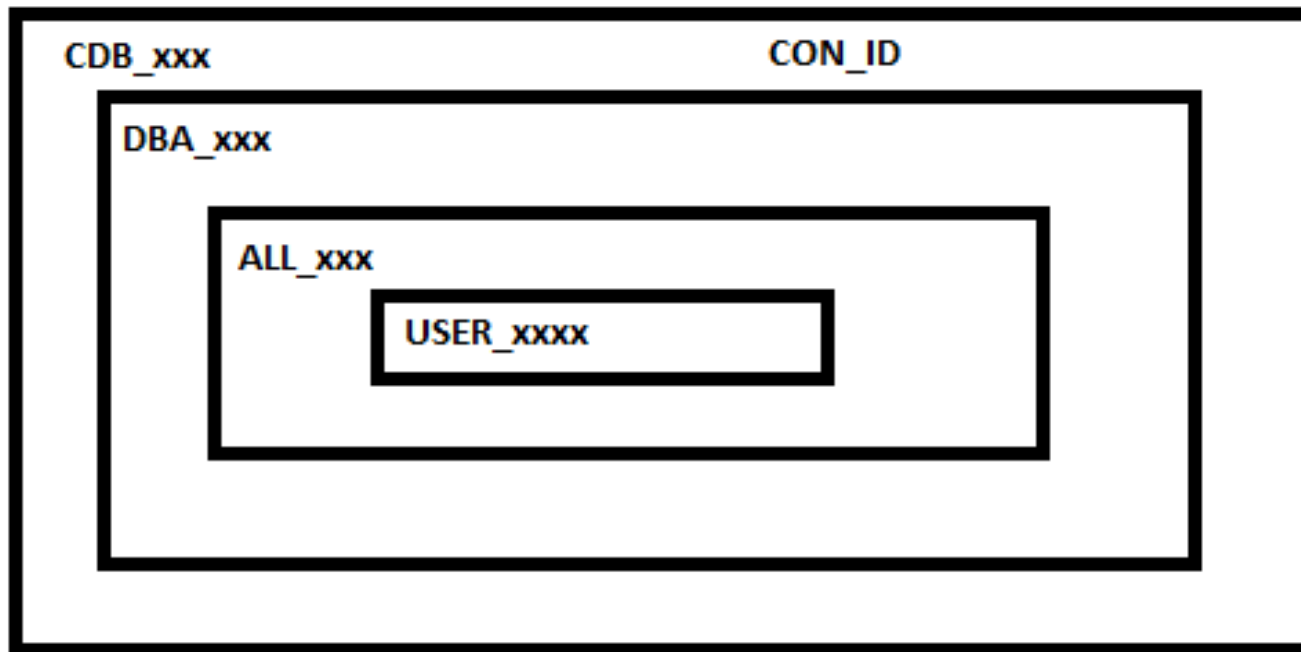
# Data Dictionary

There are three classes of dictionary views from which every user can SELECT to view information from the Dictionary:

- The USER\_\* views contain information about objects that user owns, usually because he created them.
  - Examples: USER\_TABLES, USER\_INDEXES.
- The ALL\_\* views contain information about objects that user has privileges to use.
  - These include the USER\_\* information as a subset, because one always has privileges to use the objects that he owns.
  - Examples: ALL\_TABLES, ALL\_INDEXES.
- A third class of views one can SELECT to view information from the Dictionary are normally only available to the Database Administrator:
  - The DBA\_\* tables contain information about everything in the database, no matter who owns them.
  - Examples: DBA\_TABLES, DBA\_INDEXES.

# Data Dictionary

There are three classes of dictionary views from which every user can SELECT to view information from the Dictionary:



# Data Dictionary

- Viewing Information in the Dictionary:
- The dictionary should not be modified manually.
- One can DESCRIBE and SELECT from Dictionary tables.
- Example 1: to see information about all the tables that one can use:

```
DESCRIBE ALL_TABLES
```

- Example 2: to see the name and the owner of the tables

```
SELECT table_name, owner FROM ALL_TABLES;
```

- Example 3:

```
SELECT object_type, object_name FROM USER_OBJECTS;
```

# Data Dictionary

Examples:

```
SELECT object_type, COUNT(*) FROM USER_OBJECTS  
GROUP BY object_type;
```

```
SELECT COUNT(*) FROM DICT WHERE table_name LIKE 'USER%';
```

```
SELECT * FROM DICT WHERE table_name LIKE 'USER%IND%';
```